

Beta vs Binomial Distribution

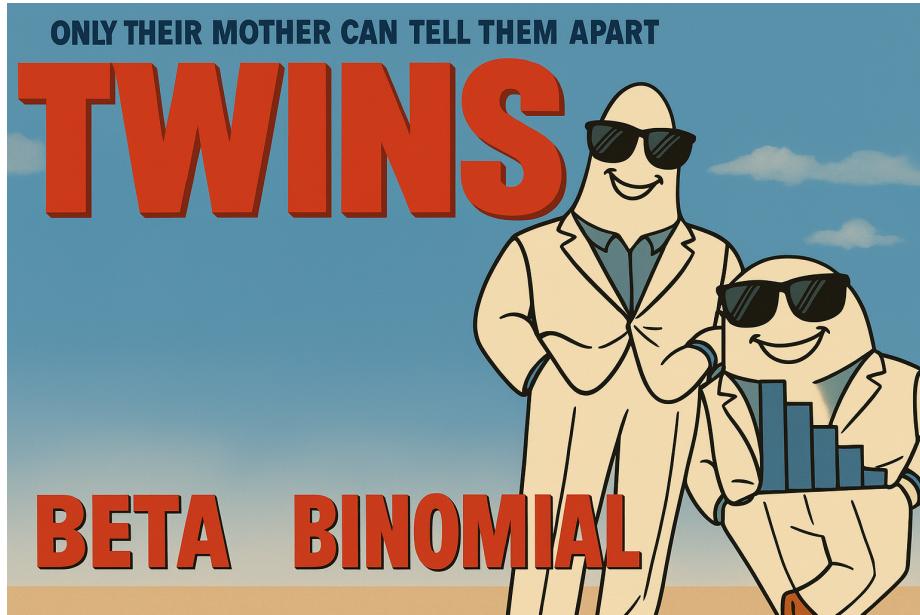
Understand the Difference Using Simulations

Russ Zaliznyak <rzaliznyak@gmail.com>

2025-12-09

Table of contents

1	Introduction	2
2	Formulation	2
3	Sim Code	3
4	Replace With Normal	9
5	Two-Proportion Test	11
5.1	Estimate P-values	11
5.2	Generate Credible Intervals	14
6	Common Pitfalls	17
7	Beta Distribution & Priors	17
8	Key Takeaways	17
9	Future Topics	18



1 Introduction

In experimentation we often estimate conversion rates or event probabilities. Understanding when to apply a Binomial model versus a Beta model is important for Significance (p-values) and Bayesian (credible intervals) analysis.

Below we present the formal definitions of our distributions, but will later run simulations to build our intuition for the Beta and Binomial Distribution.

2 Formulation

Let's spend a few minutes on the "tough" stuff: a side-by-side comparison of *Binomial Distribution* vs *Beta Distribution*. Both distributions arise from the same underlying relationship — one models outcomes given a probability, and the other models the probability itself given outcomes.

Binomial

$$P(X = k \mid n, p) = \binom{n}{k} p^k (1 - p)^{n-k}$$

$$\binom{n}{k} = \frac{n!}{k!(n - k)!}$$

i Purpose

Model the probability of observing k successes in n independent Bernoulli trials,
where each trial has success probability p .

?

Example

Given assumption that $p = 0.50$, what is the probability of generating 525 or more conversions out of 1,000 samples?

Beta

$$f(p | \alpha, \beta) = \frac{1}{B(\alpha, \beta)} p^{\alpha-1} (1-p)^{\beta-1},$$

$$B(\alpha, \beta) = \int_0^1 t^{\alpha-1} (1-t)^{\beta-1} dt = \frac{\Gamma(\alpha)\Gamma(\beta)}{\Gamma(\alpha+\beta)}.$$

i Purpose

Describe the probability of different values of the success parameter p , given successes and failures in an experiment.

?

Example

Given 525 conversions out of 1,000 samples, what are the values of p that could have generated such a test result?

Use the **Binomial** when assessing how likely your observed data are under a hypothesis; use the **Beta** when estimating the underlying conversion rate.

3 Sim Code

Now we analyze the **same** experiment with *Binomial* and *Beta* distributions.

- The *binomial* will be used to estimate the *likelihood* of our outcome (p-value).
- The *beta* will be used to estimate our *posterior distribution* of conversion rates.

```
import numpy as np
import plotly.graph_objects as go
```

```

def plot_histogram(
    data,
    bins=40,
    title=None,
    color="royalblue",
    opacity=0.7,
    normalize=False,
    show_yaxis=True,
    x_symbol="p",
    show_mean=False,
    mean_line_color="black",
    mean_line_dash="dash",
    is_rate=False,
    highlight_tails=None, # "left", "right", or "both"
    tail_alpha=0.05, # proportion per tail (e.g., 0.05 → 5%)
    tail_color="crimson",
):
    """
    Plotly histogram with optional mean line, tail highlighting, and CDF-based hover.
    """

    # Compute histogram and CDF
    hist, edges = np.histogram(
        data,
        bins=bins,
        density=normalize,
    )
    centers = 0.5 * (
        edges[1:] + edges[:-1]
    )
    cdf = np.cumsum(hist)
    cdf = (
        cdf / cdf[-1]
    ) # normalize to 1
    mean_value = np.mean(data)

    # Tail cutoffs
    left_cut = (
        np.quantile(data, tail_alpha)
        if highlight_tails
        in ["left", "both"]
        else None
    )
    right_cut = (

```

```

        np.quantile(
            data, 1 - tail_alpha
        )
        if highlight_tails
        in ["right", "both"]
        else None
    )

    # Assign colors per bar
    colors = np.full_like(
        centers, color, dtype=object
    )
    if highlight_tails:
        if left_cut is not None:
            colors[
                centers <= left_cut
            ] = tail_color
        if right_cut is not None:
            colors[
                centers >= right_cut
            ] = tail_color

    # Hover template
    if is_rate:
        hovertemplate = (
            f"P({x_symbol} < "
            + "%{customdata[0]:.2%})"
            + " = %{customdata[1]:.1%}<extra></extra>"
        )
    else:
        hovertemplate = (
            f"P({x_symbol} < "
            + "%{customdata[0]:.2f})"
            + " = %{customdata[1]:.1%}<extra></extra>"
        )

    # Single bar trace (no overlap)
    fig = go.Figure()
    fig.add_trace(
        go.Bar(
            x=centers,
            y=hist,
            marker_color=colors,
            opacity=opacity,
            customdata=np.column_stack(

```

```

        [centers, cdf]
    ),
    hovertemplate=hovertemplate,
    showlegend=False,
)
)

# Add tail annotations
annos = []

# Add mean line
if show_mean:
    annotation_text = (
        f"mean = {mean_value:.2%}"
        if is_rate
        else f"mean = {mean_value:.2f}"
    )
    fig.add_vline(
        x=mean_value,
        line_width=3,
        line_dash=mean_line_dash,
        line_color=mean_line_color,
        annotation_text=annotation_text,
        annotation_position="top",
    )

# Layout styling
fig.update_layout(
    title=title,
    xaxis_title=f"{x_symbol}",
    template="simple_white",
    bargap=0.05,
    height=320,
    yaxis=dict(
        visible=show_yaxis,
        showticklabels=False,
        title="",
    ),
    xaxis=dict(
        tickformat=".0%"
        if is_rate
        else None,
        showline=True,
        showgrid=False,
        zeroline=False,
    )
)

```

```

        ),
        margin=dict(
            t=30, b=30, l=30, r=30
        ),
    )

    return fig
}

import pandas as pd
from IPython.display import display

conversions = 525
trials = 1000
null_rate = 0.50

df = pd.DataFrame([
    "Trials": f"{trials:,}",
    "Conversions": conversions,
    "Null Hypothesis": f"{null_rate:.1%}"
])

df = df.round(2)

```

Table 1

Trials	Conversions	Null Hypothesis
1,000	525	50.0%

```

from numpy.random import binomial, seed
from numpy import (
    where,
    mean,
    percentile,
)
import plotly.graph_objects as go

NUMBER_SIMULATIONS = int(5e6)
seed(5)

# Create 5 Million Simulations each with
# n = trials and p = null_rate
control_simulated_experiments = (
    binomial(
        n=trials,

```

```

        p=null_rate,
        size=NUMBER_SIMULATIONS,
    )
)
p_value = mean(
    where(
        control_simulated_experiments
        >= conversions,
        1,
        0,
    )
)
fig = plot_histogram(
    control_simulated_experiments,
    x_symbol="conversions",
    highlight_tails="right",
    tail_alpha=p_value,
)

```

`fig.show()`

Unable to display output for mime type(s): text/html

Unable to display output for mime type(s): text/html

p-value = 0.061

The **p-value** is the likelihood of seeing our experiment data (or more extreme); assuming the Null Hypothesis is true.

```

# Observed data
from numpy.random import beta

failures = trials - conversions
tail_alpha = 0.025

beta_control_samples = beta(
    conversions,
    failures,
    size=NUMBER_SIMULATIONS,
)

fig = plot_histogram(
    beta_control_samples,
    x_symbol="Possible Conversion Rate",
    is_rate=True,

```

```

        highlight_tails="both",
        tail_alpha=tail_alpha,
        color="crimson",
        tail_color="royalblue",
    )

lower_bound = percentile(
    beta_control_samples, tail_alpha * 100
)
upper_bound = percentile(
    beta_control_samples,
    100 - tail_alpha * 100,
)

fig.show()

```

Unable to display output for mime type(s): text/html

95% Credible Interval: (49.40% , 55.59%)

The **95% credible interval** from the Beta distribution represents where the true rate most likely lies given our data.

4 Replace With Normal

Both distributions can be well approximated with the *Normal Distribution* because of CLT. Note: The Normal approximation breaks down when n is small or when p is near 0 or 1, where the distributions become skewed.

```

from numpy.random import normal
from numpy import where, mean
import plotly.graph_objects as go

# Create 1 Million Simulations each with
# n = trials and p = null_rate
mu = trials * null_rate
sigma = (
    trials * null_rate * (1 - null_rate)
) ** 0.5

control_simulated_experiments = (
    np.random.normal(
        loc=mu,
        scale=sigma,
        size=NUMBER_SIMULATIONS,

```

```

        )
    )
p_value = mean(
    where(
        control_simulated_experiments
        >= conversions,
        1,
        0,
    )
)
fig = plot_histogram(
    control_simulated_experiments,
    x_symbol="conversions",
    highlight_tails="right",
    tail_alpha=p_value,
)
fig.show()

```

Unable to display output for mime type(s): text/html
p-value = 0.0568

$$\text{Binomial}(n, p) \approx \mathcal{N}(np, np(1-p))$$

```

# Observed data
alpha = conversions
beta_param = failures
tail_alpha = 0.025

# --- Normal approximation parameters ---
mu = alpha / (alpha + beta_param)
sigma = (
    alpha
    * beta_param
    / (
        (alpha + beta_param) ** 2
        * (alpha + beta_param + 1)
    )
) ** 0.5

# Draw samples from the Normal approximation
beta_control_samples = normal(
    loc=mu,
    scale=sigma,

```

```

        size=NUMBER_SIMULATIONS,
    )

fig = plot_histogram(
    beta_control_samples,
    x_symbol="Possible Conversion Rate",
    is_rate=True,
    highlight_tails="both",
    tail_alpha=tail_alpha,
    color="crimson",
    tail_color="royalblue",
)

lower_bound = percentile(
    beta_control_samples,
    tail_alpha * 100,
)
upper_bound = percentile(
    beta_control_samples,
    100 - tail_alpha * 100,
)

fig.show()

```

Unable to display output for mime type(s): text/html

95% Credible Interval: (49.41% , 55.59%)

$$\text{Beta}(\alpha, \beta) \approx \mathcal{N}\left(\mu = \frac{\alpha}{\alpha + \beta}, \sigma^2 = \frac{\alpha\beta}{(\alpha + \beta)^2(\alpha + \beta + 1)}\right).$$

5 Two-Proportion Test

To perform an analogous analysis for the difference of two proportions, we model the difference in `# of conversions` and `delta-to-control` using the difference of our simulations for *treatment* and *control* groups.

5.1 Estimate P-values

```

import pandas as pd
from IPython.display import display

```

```

control_events = 500
treatment_events = 535
delta_conversions = treatment_events - control_events
trials = 1000
null_rate = 0

df = pd.DataFrame([{
    "control_trials": f"{trials:,}",
    "control_events": control_events,
    "treatment_trials": f"{trials:,}",
    "treatment_events": treatment_events,
    "Null Hypothesis": f"{0:.1%}"
}])

df = df.round(2)

display(df.style.hide(axis="index"))

```

Table 2

control_trials	control_events	treatment_trials	treatment_events	Null Hypothesis
1,000	500	1,000	535	0.0%

```

blended_rate = (
    control_events + treatment_events
) / (2 * trials)
control_simulated_experiments = (
    binomial(
        n=trials,
        p=blended_rate,
        size=NUMBER_SIMULATIONS,
    )
)
treatment_simulated_experiments = (
    binomial(
        n=trials,
        p=blended_rate,
        size=NUMBER_SIMULATIONS,
    )
)
delta_simulated_experiments = (
    treatment_simulated_experiments
    - control_simulated_experiments
)

```

```

)
p_value = mean(
    where(
        delta_simulated_experiments
        >= delta_conversions,
        1,
        0,
    )
)
fig = plot_histogram(
    delta_simulated_experiments,
    x_symbol="Difference of Conversions",
    highlight_tails="right",
    tail_alpha=p_value,
)

```

fig.show()

Unable to display output for mime type(s): text/html

p-value = 0.0611

```

control_se = (
    trials
    * blended_rate
    * (1 - blended_rate)
) ** 0.5
treatment_se = control_se
delta_se = (
    control_se**2 + treatment_se**2
) ** 0.5

delta_simulated_experiments = normal(
    loc=0,
    scale=delta_se,
    size=NUMBER_SIMULATIONS,
)
p_value = mean(
    where(
        delta_simulated_experiments
        >= delta_conversions,
        1,
        0,
    )
)

```

```

fig = plot_histogram(
    delta_simulated_experiments,
    x_symbol="Difference of Conversions",
    highlight_tails="right",
    tail_alpha=p_value,
)

fig.show()

```

Unable to display output for mime type(s): text/html
p-value = 0.0586

5.2 Generate Credible Intervals

```

# Observed data
from numpy.random import beta

control_failures = (
    trials - control_events
)
treatment_failures = (
    trials - treatment_events
)
tail_alpha = 0.025

beta_control_samples = beta(
    control_events,
    control_failures,
    size=NUMBER_SIMULATIONS,
)
beta_treatment_samples = beta(
    treatment_events,
    treatment_failures,
    size=NUMBER_SIMULATIONS,
)
beta_delta_samples = (
    beta_treatment_samples
    - beta_control_samples
)
fig = plot_histogram(
    beta_delta_samples,
    x_symbol="Possible Difference of Conversion Rate",
    is_rate=True,
)

```

```

        highlight_tails="both",
        tail_alpha=tail_alpha,
        color="crimson",
        tail_color="royalblue",
    )

lower_bound = percentile(
    beta_delta_samples, tail_alpha * 100
)
upper_bound = percentile(
    beta_delta_samples,
    100 - tail_alpha * 100,
)

fig.show()

```

Unable to display output for mime type(s): text/html

95% Credible Interval: (-0.88% , 7.86%)

```

tail_alpha = 0.025

# --- Normal approximation parameters ---
control_mu = control_events / trials
control_sigma = (
    control_events
    * control_failures
    / (
        (
            control_events
            + control_failures
        )
        **
        2
    *
    (
        control_events
        + control_failures
        +
        1
    )
)
) **
0.5

treatment_mu = treatment_events / trials
treatment_sigma = (
    treatment_events
    * treatment_failures
)
```

```

    / (
      (
        treatment_events
        + treatment_failures
      )
      ** 2
      * (
        treatment_events
        + treatment_failures
        + 1
      )
    )
) ** 0.5

delta_mu = treatment_mu - control_mu
delta_se = (
    control_sigma**2
    + treatment_sigma**2
) ** 0.5

# Draw samples from the Normal approximation
beta_delta_samples = normal(
    loc=delta_mu,
    scale=delta_se,
    size=NUMBER_SIMULATIONS,
)

fig = plot_histogram(
    beta_delta_samples,
    x_symbol="Possible Difference of Conversion Rate",
    is_rate=True,
    highlight_tails="both",
    tail_alpha=tail_alpha,
    color="crimson",
    tail_color="royalblue",
)
lower_bound = percentile(
    beta_delta_samples, tail_alpha * 100
)
upper_bound = percentile(
    beta_delta_samples,
    100 - tail_alpha * 100,
)

```

```

)
fig.show()

Unable to display output for mime type(s): text/html
95% Credible Interval: ( -0.88% , 7.88% )

```

6 Common Pitfalls

1. Using the normal approximation when $np < 5$ or $n(1-p) < 5$ can distort tail probabilities.
2. Confusing p -values (likelihoods of data under H_0) with posterior probabilities about p .
3. Ignoring the $+1$ in Beta(1, 1) priors isn't catastrophic, but slightly understates uncertainty when n is small.

7 Beta Distribution & Priors

The **Beta distribution** naturally supports *updating beliefs* as new data arrive.

$$\text{Prior: } \text{Beta}(\alpha_0, \beta_0) \xrightarrow{\text{observe } k, n} \text{Posterior: } \text{Beta}(\alpha_0 + k, \beta_0 + n - k)$$

This property allows the model to evolve over time — today's posterior becomes tomorrow's prior,
making the Beta approach ideal for iterative experiments or streaming data.

8 Key Takeaways

The Binomial distribution quantifies how surprising the observed data are under a fixed p_0 , while the Beta distribution estimates which values of p are most plausible given the outcomes.

Under large-sample conditions, both are well-approximated by the Normal distribution, though this breaks down when n is small or p is near 0 or 1.

The Beta model's ability to update priors as new data arrive makes it especially useful for iterative experiments or streaming environments, where today's posterior becomes tomorrow's prior.

9 Future Topics

Request future training sessions: Bayesian Priors, Analyzing Continuous Data, SRM, CUPED, and other AB Testing topics.